

PROFILE

AI carries risk. Code bounds it. The standard answers for what remains.

Programmer. Eight years of self-directed study and seven months of non-stop building: AI pipelines that run on real data, with measured numbers, written limitations and public code. I work from one simple idea: every AI application has five steps and only one invokes a model; the other four are code. When those are done properly, what is left in the middle fits in one call, and the result is cheaper, stronger and auditable. I have worked with automotive and food companies where regulation rules and a silent error is expensive.

PUBLIC PROJECTS

CTSM

[epifaneia/ctsm](#) · 2026

Audits the requirements of a customer specification (a Polarion ReqIF) against the supplier's manuals and returns the same file with the evidence and its exact page. The model may only cite fragments from a closed list; code verifies every quote against the source and derives the page from the offset. **9/9 against ground truth · 8/9 with a local 14B model on a 6 GB GPU and zero false "compliant" · round-trip diff tested against 5 sabotages.**

reqif-extractor

[epifaneia/reqif-extractor](#) · 2026

From a specification PDF to a ReqIF that Polarion imports on the first attempt: verbatim identifiers, document hierarchy, page order. Deterministic identifier policy and type reconciliation against the source. **167 of 170 model-invented identifiers caught · 1/20 of the vision cost by routing only tables and figures · 12/12 real documents structurally validated.**

toolkit

[epifaneia/toolkit](#) · 2026

The reusable part of every pipeline: one tool per design step, generic guardrails and three integration wires to the company's systems (identity, event ledger, sign-off), with tests and a documented integration for Entra ID, SIEM and PKI.

viabilidad-ia

[epifaneia/viabilidad-ia](#) · 2026

From a meeting form to an AI feasibility study in PDF: research with cited sources, structured extraction and a ranking computed by code ($\sqrt{\text{ROI} \times \text{feasibility}}$), not by the model. Form, FastAPI server, consultant panel, QA warnings and human review before sending.

Mildew Detection in Cherry Leaves

[cefeidas/Mildew-Detection-in-Cherry-Leaves](#) · 2024

Image classifier that tells healthy cherry leaves from those with powdery mildew: a convolutional network in TensorFlow/Keras trained on **5,972 class-balanced images**, 70/10/20 split, **99.8% accuracy on the test set**. Streamlit dashboard deployed on Heroku.

EXPERIENCE

AI Solutions Designer

March 2026 – present

Epifaneia · Seville, remote

Design and delivery of AI pipelines with a chain of custody for industry and regulated processes: the three projects above, a mobile voice-conversation app with a model per stage, and proofs of concept for automotive and food companies.

Automation programmer and sales agent

March – December 2024

Konecta · Seville

Automated the team's reporting with Python (Pandas, openpyxl) over Metabase exports, in iterations reviewed by the data analysis team; alongside sales activity.

Risk analysis, quality and training

2017 – 2023

TaskUs, Prepaid Financial Services, Media Interactiva, Field Management · Ireland and Spain

Financial risk analysis, quality auditing and agent training at TaskUs, including an AI model training programme; software QA at Media Interactiva. Six years working in English.

EDUCATION

Diploma in Full Stack Software Development

2022 – 2024

Code Institute · university credit-rated (University of the West of Scotland)

The only university credit-rated online coding bootcamp in the UK and Europe. Five portfolio projects at [github.com/cefeidas](#): an accessible website (HTML, CSS), a JavaScript game, a Python game on the Google Sheets API, a Django library with authentication and reviews, and the convolutional network above.

Continuous self-directed study in AI, software architecture and regulation

2018 – present

University studies in Philosophy and in Physics

University of Seville



Daniel Martín

AI Engineer

AI solutions for regulated processes

📍 Seville, Spain · remote

📞 +34 644 940 525

✉ danielmartin@epifaneia.dev

🌐 epifaneia.dev

🐙 github.com/epifaneia

🌐 linkedin.com/in/epifaneia

🔗 AVAILABLE IMMEDIATELY

LANGUAGES

Spanish native

English C1 · six years working in English

SKILLS

Applied AI

LLM pipelines with structured output, dynamic enums and deterministic verification. Local models (Ollama, Qwen 7B/14B) and cloud (Gemini) under one contract. BM25 retrieval. Calibration and measurement against ground truth. Agents with CrewAI (Flows and Crews with RAG) when the problem calls for them.

Machine learning and data

Convolutional networks with TensorFlow/Keras: data preparation and class balancing, train/validation/test split, evaluation. Pandas, NumPy, Matplotlib, Jupyter, Streamlit. PDF (PyMuPDF, Docling, gmf), XML and ReqIF (lxml), SQLite.

Code

Python (FastAPI, Django, lxml), JavaScript and TypeScript (React Native, Expo), HTML and CSS. Git, Docker, Heroku, tests and coverage. AI-assisted development with my own review, verification and guardrails.

Regulation

EU AI Act, ISO/IEC 42001, ISO/IEC 27001, GDPR, TISAX. ISO 26262 and ISO/SAE 21434 context. Polarion.

INTERESTS

Skating · Philosophy · Role-playing